

Knowledge Embedding

Wantee Wang

2015-03-24 10:05:58 +0800

Contents

1	Word2Vec	1
1.1	Continuous Bag-of-Words(CBOW) Model	3
1.2	Continuous Skip-gram Model	3
2	Paragraph to Vector	4
3	Graph to Vector	5
3.1	LINE with First-order Proximity	5
3.2	LINE with Second-order Proximity	6

Neural Networks are applied in many fields of Machine Learning. Due to their natural property, it is suitable to use them to estimate the *distribution representations* from some knowledge source. This kind of representation may be referred as *knowledge embedding*.

[Embedding](#), mathematically, denotes one instance of some mathematical structure contained within another instance. In Machine Learning context, it often means to transfer a sparse coding of an instance to a more dense coding, e.g., a 1-of-V coding to a low dimension vector.

1 Word2Vec

The first successful application of NN-based embedding is [word2vec](#), which converts a word to a continuous vector.

The inspiration [1] of word2vec is from the observation that neural network language model can be successfully trained in two steps: first, continuous word vectors are learned using simple model, and then the N-gram NNLM is trained

on top of these distributed representations of words. Then the authors spend lots of work on learning word vectors. Thus the word-to-vector procedure can be seen as a feature extraction step for language modelling.

To show the innovation of word2vec, we first consider the traditional N-gram NNLM. The architecture of NNLM [2] is shown in following figure (Figure 1),

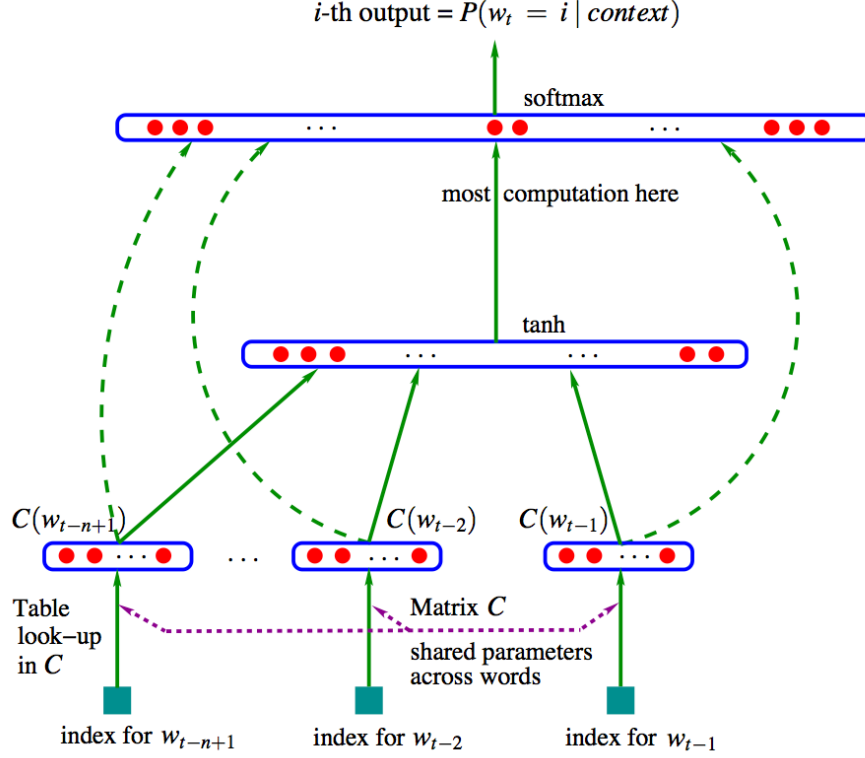


Figure 1: Architecture of NNLM

The NNLM consists of input, projection, hidden and output layers. At the input layer, N previous words are encoded using 1-of- V coding, where V is size of the vocabulary. The input layer is then projected to a projection layer that has dimensionality $N \times D$. Next, projection layer is connected to the hidden layer, whose size is H , finally, we reach the output layer, which is of size V , through hidden layer.

Therefore, the computational complexity per each training example is

$$Q = N \times D + N \times D \times H + H \times V$$

Where the domination term is $H \times V$. However several output layer optimisations method can avoid it, such as *Hierarchical Softmax* which reduce the output term to $H \times \log_2(V)$. Then the domination term becomes $N \times D \times H$.

It can be seen from above discussion the most complexity is caused by the non-linear hidden layer in the model (if we check out the RNN LM model, same conclusion can be derived, i.e., the hidden layer occupied most computational resources).

However, if the goal is only to extract word embeddings, we can sacrifice some precision of NN models. This leads to the key point of speedup of word2vec, that is by **removing the non-linear hidden layer**, and used a log-linear model directly.

1.1 Continuous Bag-of-Words(CBOW) Model

There are two type of log-linear models, first one is the bag-of-words model. As the name showed, it don't consider the order of the words in history. The architecture is the same with NNLM, except that the hidden layer is removed and the projection layer is shared for all words (not just the projection matrix), as shown in the figure (Figure 2),

Another different point with traditional models is word2vec using both pre- and post-contexts of the predicting word. The training complexity is then

$$Q = N \times D + D \times \log_2(V)$$

1.2 Continuous Skip-gram Model

The second model is similar to CBOW, but instead of predicting the current word based on the context, it tries to maximise classification of a word based on another word in the same sentence. More precisely, it uses each current word as an input to a log-linear classifier with continuous projection layer, and predict words within a certain range before and after the current word. The architecture is shown in the following figure (Figure 3),

The complexity of Skip-gram model is

$$Q = 2 \times C \times (D + D \times \log_2(V))$$

where C is the maximum distance of the words. Thus, for a particular C , for each training word it will select randomly a number $R \in [1, C]$, and then use R words from history and R words from the future of the current word as correct labels.

For both CBOW and Skip-gram models, after training, the corresponding column of the projection matrix is taken out to be the vectors of one word.

The objective function of Skip-gram models is to maximise the average log probability

$$\mathcal{O} = \frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_{t+j} | w_t) \quad (1)$$

where the probability $p(w_{t+j} | w_t)$ is defined by softmax function,

$$p(w_O | w_I) = \frac{e^{\mathbf{v}'_{w_O} \mathbf{v}_{w_I}}}{\sum_{w=1}^V e^{\mathbf{v}'_w \mathbf{v}_{w_I}}} \quad (2)$$

where \mathbf{v}_w and \mathbf{v}'_w are the “input” and “output” vector representation of w , which are the corresponding column of projection matrix and the weight matrix between projection and output layer.

After above work, [3] proposed some method to further speedup the training, such as hierarchical softmax and negative sampling, and then extended words to phrases. We won’t go deep into these in this post.

2 Paragraph to Vector

Inspired by word2vec, [4] extends it to transform a variable-length of text to a *paragraph vector*. The architecture is similar with CBOW, as shown in the following figure (Figure 4), except that the additional paragraph token in input layer, which is mapped to a vector via matrix D .

The paragraph token can be thought of as another word. It acts as a memory that remembers what is missing from the current context – or the topic of the paragraph. The contexts are fixed-length and sampled from a sliding window over the paragraph. The paragraph vector is shared across all contexts generated from the same paragraph but not across paragraphs. The word vector matrix W , how-ever, is shared across paragraphs. At every step of stochastic gradient descent, one can sample a fixed-length context from a random paragraph, then use them to train the network. At prediction time, one needs to perform an inference step to compute the paragraph vector for a new paragraph. This is also obtained by gradient descent. In this step, one can add one column in D and gradient descending on D while holding W fixed.

After being trained, the paragraph vectors can also be used as features for the paragraph and be feed directly to other classifiers.

The most important advantages of paragraph vectors is that they are learned from unlabelled data. Besides, paragraph vectors also address some of the key weaknesses of bag-of-words models. First, they inherit an important property of the word vectors: the semantics of the words. The second advantage of the paragraph vectors is that they take into consideration the word order, at least in a small context.

3 Graph to Vector

[5] further extends the embedding idea to general information networks, more specifically, it transfer the vertices in a graph to vectors.

The goal is to use a low-dimensional vector to represent a vertex in the graph, while preserving both local and global structure informations. To derive the model, they first formally defined the local and global similarity of vertices,

The local similarity is defined by **First-order Proximity**, which is the weight w_{uv} on the edge that connected vertex u and vertex v .

The **Second-order Proximity** of a pair (u, v) is the similarity between their neighbourhood network structure. Mathematically, let $\mathbf{p}_u = (w_{u,1}, \dots, w_{u,|V|})$ denotes the first-order proximity of u with all the other vertices, then the second-order proximity between u and v is determined by the similarity between \mathbf{p}_u and \mathbf{p}_v . The second-order proximity assumes that vertices sharing many connections to other vertices are similar to each other.

Thus, the graph embedding problem becomes that to convert a vertex to vector, preserving the first- and second-order proximities.

Their model is called *Large-scale Information Network Embedding (LINE)*. A graph is denoted by $G = (V, E)$, where V is the set of vertices and E is the set of edges.

3.1 LINE with First-order Proximity

To model the first-order proximity, for each undirected edge (i, j) , define the joint probability between vertex v_i and v_j as

$$p_1(v_i, v_j) = \frac{1}{1 + e^{-\mathbf{u}_i^\top \mathbf{u}_j}} \quad (3)$$

where $\mathbf{u}_i \in \mathbb{R}^d$ is the d -dimensional vector of vertex v_i .

Note that the empirical probability of $p_1(\cdot, \cdot)$ can be defined as

$$\hat{p}_1(v_i, v_j) = \frac{w_{ij}}{\sum_{(i,j) \in E} w_{ij}} \quad (4)$$

Thus the objective function of first-order proximity, is to minimise the following function

$$\mathcal{O}_1 = d(\hat{p}_1(\cdot, \cdot), p_1(\cdot, \cdot)) \quad (5)$$

where $d(\cdot, \cdot)$ is the distance between two distributions. Replacing it with the KL-divergence and omitting some constants¹,

$$\mathcal{O}_1 = - \sum_{(i,j) \in E} w_{ij} \log p_1(v_i, v_j) \quad (6)$$

Note that the first-order proximity is only applicable for undirected graphs, not for directed ones.

3.2 LINE with Second-order Proximity

The second-order proximity is applicable for both directed and undirected graphs. Given a network, without loss of generality, we assume it is directed.

Each vertex can be treated as a specific “context” and vertices with similar distributions over the “contexts” are assumed to be similar. Therefore, each vertex plays two roles: the vertex itself and a specific “context” of other vertices.

Let \mathbf{u}_i be the representation of v_i when it is treated as a vertex, while \mathbf{u}'_i is the representation of v_i when it is treated as a specific “context”. For each directed edge (i, j) , first define the probability of “context” v_j generated by vertex v_i as

$$p_2(v_j | v_i) = \frac{e^{\mathbf{u}'_j \top \mathbf{u}_i}}{\sum_{k=1}^{|V|} e^{\mathbf{u}'_k \top \mathbf{u}_i}} \quad (7)$$

where $|V|$ is the number of vertices or “contexts”.

And the empirical probability is

$$\hat{p}_2(v_j | v_i) = \frac{w_{ij}}{d_i} \quad (8)$$

where d_i is the out-degree of vertex i , i.e. $d_i = \sum_{k \in N(i)} w_{ik}$, where $N(i)$ is the set of out-neighbours of v_i .

Similarly, the objective of second-order proximity is

¹The relationship between cross entropy and KL distance is $H(\hat{p}, p) = H(\hat{p}) + D_{KL}(\hat{p} | p)$. For the purpose of optimising the objective, $H(\hat{p})$ is constant, thus the objective is same as cross entropy.

$$\mathcal{O}_2 = \sum_{i \in V} \lambda_i d(\hat{p}_2(\cdot | v_i), p_2(\cdot | v_i)) \quad (9)$$

For simplicity, set $\lambda_i = d_i$, then replacing $d(\cdot, \cdot)$ with the KL-divergence and omitting some constants,

$$\mathcal{O}_2 = - \sum_{(i,j) \in E} w_{ij} \log p_2(v_j | v_i) \quad (10)$$

References

- [1] Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. [Efficient estimation of word representations in vector space](#). *arXiv preprint arXiv:1301.3781*, 2013.
- [2] Yoshua Bengio, Réjean Ducharme, Pascal Vincent, and Christian Janvin. [A neural probabilistic language model](#). *The Journal of Machine Learning Research*, 3:1137–1155, 2003.
- [3] Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg S Corrado, and Jeff Dean. [Distributed representations of words and phrases and their compositionality](#). In *Advances in Neural Information Processing Systems*, pages 3111–3119, 2013.
- [4] Quoc V Le and Tomas Mikolov. [Distributed representations of sentences and documents](#). *arXiv preprint arXiv:1405.4053*, 2014.
- [5] Jian Tang, Meng Qu, Mingzhe Wang, Ming Zhang, Jun Yan, and Qiaozhu Mei. [LINE: Large-scale Information Network Embedding](#). *arXiv preprint arXiv:1503.03578*, 2015.

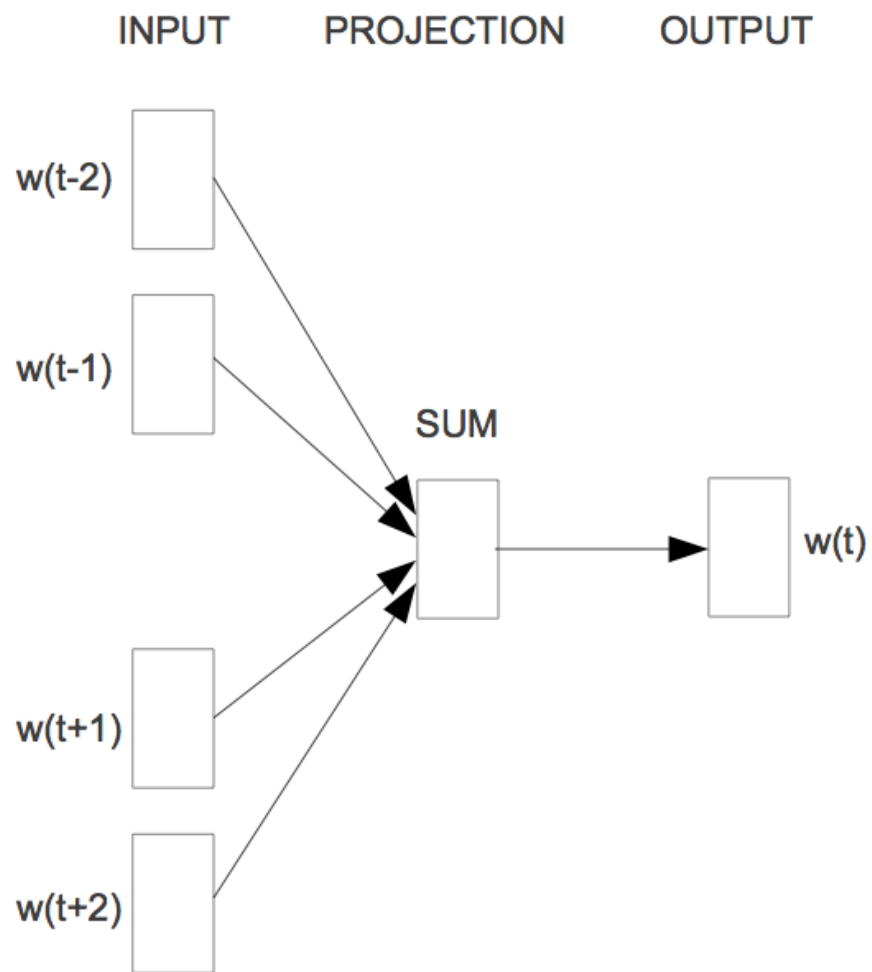


Figure 2: Architecture of CBOW

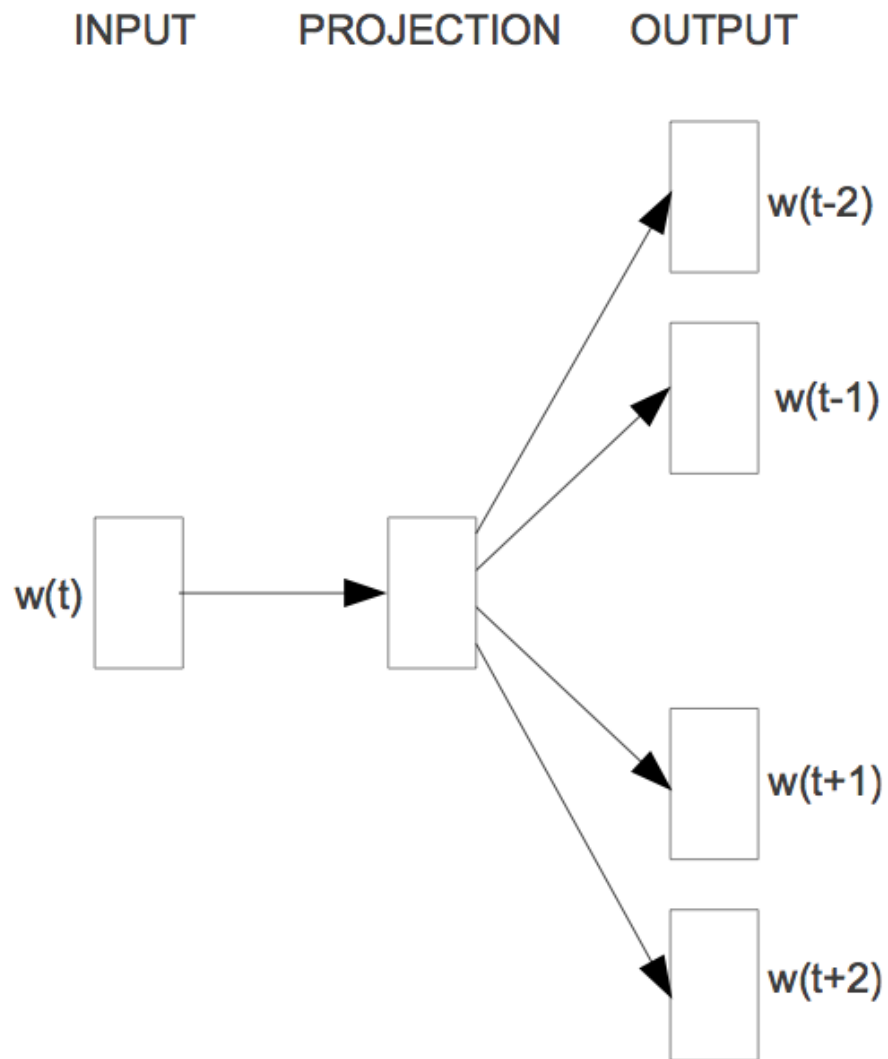


Figure 3: Architecture of Skip-gram

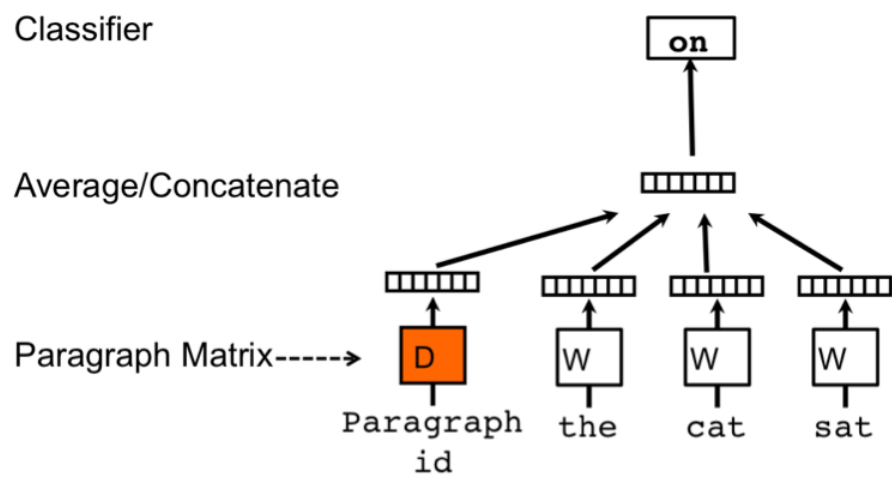


Figure 4: Architecture of Paragraph Vector